

STEVENS INSTITUTE OF TECHNOLOGY

BIA-678: Big Data Technologies

Scalable Big Data Pipeline for Financial Fraud Detection

Pratik Kale (20036601), Pooja Pande (20035197), Vidhi Palan (20031836)

May 11, 2025

Table of Contents

1. INTRODUCTION	3
2. PROBLEM STATEMENT	4
2.1 PROJECT OBJECTIVE:	4
3. DATASET OVERVIEW & DATA DICTIONARY	5
3.1 DATA DICTIONARY	5
4. DATA PRE-PROCESSING & FEATURE ENGINEERING	6
4.1 DATA CLEANING	7
4.2 FEATURE CONSTRUCTION	8
4.3 ENCODING & SCALING	9
5. EXPLORATORY DATA ANALYSIS (EDA).....	10
5.1 CLASS DISTRIBUTION & RARITY	10
5.2 TRANSACTION AMOUNT CHARACTERISTICS.....	11
5.3 TYPE-WISE FRAUD INCIDENCE.....	11
5.4 BALANCE TRANSITION PATTERNS	11
5.5 FEATURE CORRELATION ANALYSIS	11
6. SOLUTION APPROACH & MODELING.....	12
6.1 MODEL SELECTION CRITERIA	12
6.2 TRAINING & EVALUATION PIPELINE	13
7. PERFORMANCE METRICS & RESULTS.....	14
7.1 CONFUSION MATRIX ANALYSIS	15
7.2 ROC CURVE & THRESHOLD SELECTION	15
7.3 BUSINESS IMPACT SIMULATION	15
7.4 INSIGHTS & NEXT STEPS	16
8. HYPERPARAMETER TUNING & OPTIMIZATION	16
8.1 OPTIMIZATION FRAMEWORK	16
8.2 SEARCH SPACE DEFINITION.....	17
8.3 TRIAL EXECUTION & PRUNING	17
8.4 BEST-FOUND CONFIGURATION	17
8.5 INSIGHTS & RECOMMENDATIONS.....	17
9. SCALABILITY: SCALE UP & SCALE OUT.....	18
9.1 SCALE UP: DATA VOLUME.....	18
9.2 SCALE OUT: CLUSTER SIZE	18
9.3 RECOMMENDATIONS FOR PRODUCTION	19
10. APPLICATION UI & DEPLOYMENT	19
11. CONCLUSION & FUTURE ENHANCEMENTS	20
12. REFERENCES.....	21

1. Introduction

Digital payment platforms have revolutionized commerce by enabling instant, borderless transactions between individuals and businesses worldwide. However, this convenience and speed have also created fertile ground for sophisticated and adaptive fraudsters. As transaction volumes escalate—often into the tens of thousands per second—manual review becomes not only impractical but nearly impossible. Even rule-based systems, which rely on fixed thresholds or static heuristics, struggle to keep pace with evolving attack patterns and can generate overwhelming numbers of false alarms, draining resources and eroding customer trust.

In response, our project leverages the power of big-data frameworks—specifically Apache Spark and the PySpark API—to construct an end-to-end fraud detection pipeline that addresses these challenges head-on. By processing raw transaction streams in parallel across distributed clusters, we achieve both the low latency required for near real-time scoring and the throughput needed to analyze millions of events per hour. Crucially, we pair this scalable infrastructure with advanced behavioral feature engineering—such as balance deltas, ratio metrics, and anomaly flags—that distills complex transactional histories into concise, predictive inputs. These inputs feed into ensemble machine-learning models capable of capturing non-linear relationships and subtle patterns indicative of fraud.

This report details our methodology, beginning with data ingestion and cleaning, progressing through feature construction and exploratory analysis, and culminating in model training, evaluation, and deployment. We emphasize three pillars throughout:

- **Accuracy & Robustness:** Employing techniques like SMOTE for class imbalance and Optuna for hyperparameter tuning ensures our model maintains high recall on fraud cases without sacrificing precision.
- **Scalability & Efficiency:** A distributed Spark-based pipeline allows rapid experimentation and seamless scaling—from a single node to multi-node clusters—across terabytes of data.
- **Production Readiness:** We demonstrate how to containerize the entire workflow, integrate CI/CD pipelines, and deploy to Kubernetes-managed environments, enabling continuous monitoring and retraining.

By the end of this report, the reader will understand not only the theoretical underpinnings of our fraud detection approach but also the practical considerations for implementation in a high-throughput, production-grade setting.

2. Problem Statement

In the era of digital banking and mobile wallets, financial fraud has evolved from isolated incidents into systematic threats that can compromise entire institutions. Fraudsters continually refine their tactics—such as account takeover, synthetic identity creation, and mule networks—leveraging automation and anonymization tools to bypass conventional defenses. These schemes inflict direct monetary losses, damage brand reputation, attract regulatory scrutiny, and erode customer confidence. Recent studies estimate global fraud losses in the hundreds of billions of dollars annually, underscoring the imperative for intelligent, scalable detection strategies.

Traditional rule-based engines suffer from several critical limitations:

- **Static Thresholds:** Fixed rules (e.g., flag any transaction above a fixed dollar amount) cannot adapt to shifting fraud patterns or regional variations in spending behavior.
- **High False-Positive Rates:** Overly rigid rules often capture legitimate outliers—such as large legitimate purchases or bulk corporate payments—forcing manual review teams to spend hours vetting false alarms.
- **Delayed Adaptation:** Updating rules requires manual analysis and deployment cycles, creating dangerous windows in which new attack vectors go undetected.

To overcome these challenges, we propose a data-driven, machine-learning-oriented solution built on distributed big-data technology. By framing fraud detection as a supervised classification task, our approach learns from historical examples to identify evolving fraud signatures. Key aspects of our problem formulation include:

1. **Feature Variability:** Instead of relying on raw amounts alone, we extract behavioral features—such as inter-transaction time gaps, balance deltas, and velocity metrics—that generalize across accounts and geographies.
2. **Imbalance Handling:** With fraudulent cases representing only a fraction of overall volume, specialized techniques (e.g., SMOTE, class-weighted algorithms) are crucial to prevent model bias toward the majority class.
3. **Real-Time Scoring:** Latency must be low (sub-second) to allow fraud-blocking actions (e.g., OTP triggers, transaction holds) without disrupting customer experience.
4. **Regulatory Compliance:** Models must provide explainable signals for auditing (e.g., feature importance, rule extraction) to satisfy regulators and streamline dispute resolution.

2.1 Project Objective:

- **Accuracy & Precision:** Achieve a recall of at least 75% on fraud cases while keeping false-positive rates below 5%, balancing detection power and operational cost.

- **Scalability & Throughput:** Design a pipeline capable of ingesting and processing over 10 million transactions per hour on commodity clusters, enabling growth without performance degradation.
- **Maintainability & Extensibility:** Implement a modular architecture that allows data scientists to incorporate new features, retrain models, and redeploy without interrupting production services.
- **Robust Monitoring & Drift Detection:** Establish continuous monitoring of model performance metrics and distributional shifts to trigger automated retraining workflows, ensuring sustained effectiveness.

By articulating the problem in these terms, we lay a clear foundation for a solution that not only detects today’s fraud tactics but can evolve as new threats emerge.

3. Dataset Overview & Data Dictionary

To prototype and validate our fraud detection pipeline, we employ a **synthetic mobile-money transaction dataset** designed to mimic real-world banking behavior at scale. Below are key characteristics of the dataset:

- **Volume & Coverage:** Approximately 6.3 million transactions spanning diverse transaction types (e.g., payments, transfers, cash-outs) distributed uniformly across a simulated 30-day period.
- **Synthetic Generation:** Transactions are generated via a stochastic simulator that models customer behavior, account balances, and fraud injection. This allows controlled experiments on known fraud labels while preserving realistic noise patterns.
- **Class Imbalance:** Fraudulent events constitute only ~0.12% of the dataset—reflecting true-world rarity—posing challenges for model learning without specialized oversampling or weighting strategies.

3.1 Data Dictionary

Below is a detailed description of each field, its type, and its role in subsequent analysis and modeling.

Column Name	Type	Description
step	Integer	Discrete time step (in simulation units) used to order events; dropped before modeling to avoid temporal leakage.
type	String	Transaction category: PAYMENT, TRANSFER, CASH_OUT, DEBIT, CASH_IN. Serves as a key categorical feature.
amount	Float	Monetary value of the transaction (range 0.01–100,000). Influences delta and ratio-based features.
nameOrig	String	Pseudonymized identifier of the originating account. Used for grouping behavior analysis, not in final features.

oldbalanceOrig	Float	Originator's balance immediately before the transaction. Basis for delta calculation.
newbalanceOrig	Float	Originator's balance immediately after the transaction. Basis for zero-wipeout flag and consistency checks.
nameDest	String	Pseudonymized identifier of the destination account. Included for exploratory grouping; not encoded in model.
oldbalanceDest	Float	Recipient's balance before the transaction. Used for balance delta features.
newbalanceDest	Float	Recipient's balance after the transaction. Used to validate amount consistency and for delta features.
isFraud	Integer	Label indicator (1 = Fraudulent, 0 = Genuine) used as the target variable in supervised learning.

Notes on Data Integrity & Edge Cases:

- Some records exhibit **inconsistent balances** (e.g., `oldbalanceOrig - newbalanceOrig != amount`), representing synthetic edge-case scenarios that highlight unusual behavior. We retained these to train the model on anomalies rather than discarding them outright.
- A small subset (<0.05%) of transactions contains zero or negative balances, simulating system adjustments or error recoveries; these were flagged during preprocessing and handled via capping and imputation.
- Duplicate `nameOrig` and `nameDest` pseudonyms appear across different time steps, enabling sequence- and behavior-based feature exploration (e.g., frequency of transactions per account in a rolling window).

Why Synthetic?

Using a synthetic dataset offers several advantages:

1. **Controlled Fraud Injection:** We know the exact fraud labels and types, allowing clear performance benchmarking.
2. **Privacy Preservation:** No real customer data is exposed, facilitating open-source experimentation.
3. **Scalability Testing:** Synthetic data can be scaled arbitrarily to test system limits without legal or ethical constraints.

With this comprehensive overview, we proceed to transform and engineer predictive features that distill these raw fields into high-signal inputs for our machine-learning models.

4. Data Pre-processing & Feature Engineering

To transform raw transactions into high-signal inputs for machine learning, we structured our work into three phases: data cleaning, feature construction, and encoding/scaling.

4.1 Data Cleaning

Before feature engineering, rigorous cleaning ensures data quality and integrity:

- **Initial Profiling & Schema Validation:**
 - Assessed column types, value ranges, and basic statistics (mean, median, missing counts) across all fields.
 - Validated that numeric fields (amount, oldbalanceOrg, etc.) were correctly parsed and cast from CSV strings.
- **Dropping Irrelevant & Redundant Columns:**
 - Removed simulation artifacts such as step, which only orders events without predictive power.
 - Discarded constant or near-duplicate fields (e.g., columns with >99.9% identical values) to reduce noise and memory footprint.
- **Handling Missing & Malformed Values:**
 - Conducted null checks—no nulls detected after ingestion.
 - Identified malformed balance entries (non-numeric or empty strings) and corrected via imputation using nearby time-step averages (<0.01% of rows).
- **Duplicate & Anomaly Detection:**
 - Checked for exact duplicate transaction records; eliminated <0.001% duplicates to prevent skewed learning.
 - Employed interquartile range (IQR) methods on amount and balance deltas to flag extreme outliers. Outliers were not removed but annotated for model robustness.
- **Balance Consistency Checks:**
 - Verified that $\text{oldbalanceOrg} - \text{newbalanceOrig} == \text{amount}$ and $\text{newbalanceDest} - \text{oldbalanceDest} == \text{amount}$.
 - Flagged ~0.5% of records with inconsistencies as synthetic edge cases; retained for training to help the model learn from anomalies.
- **Negative & Zero Balance Handling:**
 - Flagged transactions resulting in negative balances—often system corrections or synthetic adjustments—and retained them with a binary indicator.
 - Normalized zero-balance wipeouts by creating a zero_wipe_flag for cases where $\text{newbalanceOrig} == 0$ and $\text{oldbalanceOrg} > 0$.
- **Audit Logging & Snapshotting:**
 - Logged counts of removed columns, duplicate rows dropped, and flags set.
 - Saved original vs. cleaned snapshots for reproducibility and roll-back, storing versioned Parquet files on HDFS.

These comprehensive cleaning steps yielded a reliable foundation, reducing spurious noise and preserving meaningful anomalies for model learning.

4.2 Feature Construction

In this phase, we translate cleaned transaction records into a suite of predictive features that capture both the magnitude and the nuance of monetary flows and user behavior. Our feature set comprises five subgroups:

1. Monetary Flow Metrics

- **DeltaOrg (deltaOrg):** Computed as $\text{oldbalanceOrg} - \text{newbalanceOrig}$, this feature quantifies the exact net outflow from the originator's account. Larger deltas can indicate bulk transfers or cash-outs, while unusually small or zero deltas may signal no-fund-movement attempts (common in test or probing transactions).
- **DeltaDest (deltaDest):** Defined as $\text{newbalanceDest} - \text{oldbalanceDest}$, measuring net inflow into the recipient's account. A large positive deltaDest often accompanies fraud patterns where funds are funneled into mule accounts.

2. Relative Ratio Metrics

- **Originator Ratio (amt_to_org):** $\text{amount} / (\text{oldbalanceOrg} + 1)$. By normalizing the transaction amount against the originator's prior balance, we detect transfers that are large relative to typical account size, rather than absolute thresholds.
- **Recipient Ratio (amt_to_dest):** $\text{amount} / (\text{oldbalanceDest} + 1)$. Similar normalization for the recipient side uncovers disproportionate inflows, distinguishing between routine payroll deposits and anomalous spikes.

3. Behavioral Anomaly Flags

- **Zero-Wipeout Flag (zero_wipe_flag):** Binary indicator set to 1 when $\text{newbalanceOrig} == 0$ and $\text{oldbalanceOrg} > 0$. Fraudsters often drain accounts completely to maximize illicit gain before detection.
- **Large-Transfer Flag (large_tx_flag):** Marks transactions where amount exceeds the global mean plus three standard deviations. This thresholding captures extreme outliers that warrant heightened scrutiny.

4. Temporal and Velocity Indicators

- **Transaction Frequency:** Count of transactions per nameOrig over rolling windows (1 hour, 24 hours). Sudden bursts of activity can reflect automated attack scripts or mule-farm operations.
- **Inter-Transaction Gap:** Time elapsed since the previous transaction for the same account (step difference scaled to real units). Very short gaps may indicate robotic behavior rather than organic spending.

5. Composite & Interaction Features

- **Flow Ratio (deltaOrg / amount):** Proportion of the transfer that actually changed account state, highlighting fee-free or reversals. Values significantly below 1 may indicate simulation artifacts or reversal attempts.
- **Type–Amount Interaction:** Pairing type one-hot embeddings with amount via element-wise multiplication, enabling the model to learn type-specific thresholds (e.g., large cash-outs vs. large payments behave differently).

By combining these engineered dimensions, we create a rich, low-collinearity feature matrix that empowers tree-based models to detect subtle, non-linear fraud signatures. These features were validated through correlation analysis (all pairwise $|\rho| < 0.4$) and preliminary feature-importance screening, which highlighted deltaOrg, zero_wipe_flag, and amt_to_org among the top predictors.

4.3 Encoding & Scaling

After constructing our predictive features, we standardize and organize them into a format compatible with Spark ML pipelines. This stage ensures models train efficiently and generalize well across varied data distributions.

1. Categorical Feature Encoding

- **StringIndexer:** We first apply StringIndexer to convert the type column into a zero-based numeric index (type_idx). This step handles unseen categories in test or production data by mapping them to a reserved “unknown” index, preventing pipeline failures.
- **OneHotEncoder:** Using Spark ML’s OneHotEncoder, we transform type_idx into a sparse binary vector (type_vec). This encoding avoids imposing an ordinal relationship between transaction categories, allowing models to treat each type independently.

2. Feature Vector Assembly

- **VectorAssembler:** We then merge all numeric features (e.g., deltaOrg, amt_to_org, transaction_frequency) and the one-hot vectors (type_vec) into a single features column of type Vector. This consolidated column drives downstream learners and simplifies pipeline serialization.

3. Scaling & Normalization

- Although tree-based models are invariant to feature scaling, algorithms such as logistic regression and gradient-boosted methods benefit significantly from standardized inputs. We apply Spark ML’s StandardScaler to the assembled features vector, producing scaledFeatures with zero mean and unit variance. Using the withMean=true and withStd=true parameters:

- **Mean Centering:** Subtracts the feature's average value across the training set.
- **Unit Variance Scaling:** Divides by its standard deviation, ensuring uniform feature impact during optimization.

4. Pipeline Integration & Persistability

- All transformers (StringIndexer, OneHotEncoder, VectorAssembler, StandardScaler) are chained in a Pipeline. This encapsulated pipeline object can be saved and loaded across Spark sessions, guaranteeing consistent data transformation during training, evaluation, and production scoring.
- **Parameter Tuning:** We expose critical parameters (e.g., dropLast in OneHotEncoder, withMean flag in StandardScaler) to our hyperparameter tuning framework (Optuna), allowing automated discovery of optimal preprocessing settings.

5. Handling Skew & Sparsity

- Transaction data often exhibits high sparsity after one-hot encoding (few active types per record). By using sparse vector representations, we drastically reduce memory usage, enabling scaling to millions of rows.
- We also monitor feature distribution shifts via Spark's DataframeStatFunctions to detect when scaling parameters may need recalibration (e.g., after major market events or model drift).

With these encoding and scaling steps, our data are prepared for robust, reproducible model fitting, ensuring that both linear and non-linear learners receive appropriately structured inputs.

5. Exploratory Data Analysis (EDA)

A comprehensive EDA uncovers the underlying structures and potential signals in our transaction data. We focus on distributional properties, behavioral patterns, and feature interrelationships to guide feature selection and model design.

5.1 Class Distribution & Rarity

Fraudulent transactions account for just **0.12%** of total events, with genuine cases comprising the remaining **99.88%**. This extreme imbalance means that a naïve classifier predicting “no fraud” would achieve 99.88% accuracy yet detect zero real fraud incidents. Such a baseline underscores the necessity for specialized imbalance-handling techniques rather than relying on raw accuracy metrics.

5.2 Transaction Amount Characteristics

Contrary to the assumption that fraud correlates with large sums, we observe that fraudulent transactions tend to skew toward **smaller amounts**. This likely reflects fraudsters' attempts to evade fixed-threshold systems by dispersing illicit transfers into multiple low-value transactions. A density plot (omitted here for brevity) showed a rightward tail for genuine amounts but a pronounced left-skew for fraud.

5.3 Type-wise Fraud Incidence

Different transaction categories exhibit varying fraud rates:

In this bar chart (extracted from Slide 7), **TRANSFER** and **CASH_OUT** transactions display the highest fraud incidence, at approximately 1.5% and 1.0% respectively, compared to under 0.1% for other categories. These insights inform targeted feature interactions and risk rules for specific transaction types.

5.4 Balance Transition Patterns

By plotting originator balances before vs. after transactions, we highlight full-drain behaviors:

In this scatterplot (Slide 5), fraud cases cluster near the origin, indicating **zero-wipeout** transactions where the account is completely drained. Such patterns validate our `zero_wipe_flag` feature and motivate additional ratio-based metrics.

5.5 Feature Correlation Analysis

Understanding multicollinearity helps avoid redundant features and ensures model stability. We computed pairwise Pearson correlations across our engineered features and visualized them:

The heatmap (Slide 9) shows all pairwise correlations below **0.3** in absolute value, indicating low multicollinearity. This result confirms that our combined set of deltas, ratios, and flags provides diverse, non-redundant signals for classification.

Key Insights:

- **Behavioral Signals Over Amounts:** Ratios and delta-based features capture nuanced patterns that raw amounts miss.
- **Category-Specific Risks:** Focusing on high-risk types (**TRANSFER**, **CASH_OUT**) can improve early detection.
- **Feature Diversity:** Low correlation among engineered features supports robust learning and reduces overfitting risks.

With these EDA findings, we proceed to model selection and training in the next section.

6. Solution Approach & Modeling

Our solution approach transforms exploratory insights into a production-ready classification pipeline. We designed a modular workflow that encompasses model selection, distributed training, evaluation, and iterative refinement. Key stages include:

1. **Algorithm Benchmarking:** Compare candidate models against accuracy, runtime, and resource constraints.
2. **Pipeline Orchestration:** Leverage Spark ML Pipelines for end-to-end data flow—from feature assembly to model fitting—to ensure scalability and reproducibility.
3. **Imbalance Mitigation:** Incorporate sampling (SMOTE) and class-weighted objectives to address the 0.12% fraud prevalence.
4. **Evaluation & Monitoring:** Use cross-validated metrics (F1, ROC-AUC) and confusion matrices to balance detection power with false alarms.

6.1 Model Selection Criteria

Selecting the optimal learning algorithm requires balancing multiple dimensions:

- **Predictive Performance:** Ability to capture non-linear interactions and subtle patterns indicative of fraud.
- **Scalability & Throughput:** Training and inference must scale horizontally on Spark clusters to meet real-time scoring demands.
- **Interpretability & Auditability:** Models should provide explainable outputs (feature importances, rule approximations) to satisfy regulatory requirements.
- **Computational Efficiency:** Low training and scoring latency to minimize infrastructure cost and support online blocking actions.

Based on these criteria, we assessed the following algorithms:

Algorithm	Predictive Power	Scalability	Interpretability	Inference Latency
Logistic Regression	Moderate (linear)	Excellent	High	Very Low
Random Forest	High (non-linear)	Good (via Spark)	Medium (feature importance)	Low
Isolation Forest	Medium (anomaly detection)	Excellent	Low (unsupervised)	Very Low

Logistic Regression served as our baseline due to its simplicity and transparency. It trains quickly on large datasets and yields coefficients that directly inform risk rules, but its linear decision boundary limited recall under complex fraud scenarios.

Random Forest offered the best compromise, capturing multi-dimensional feature interactions and achieving superior F1-scores in preliminary tests. While training hundreds of trees incurs higher compute cost, Spark's parallelism and tree-parallel training model mitigated overhead.

Isolation Forest, an unsupervised approach, excels at detecting outliers without labelled data but underperformed in precision when compared to supervised methods, given the severe class imbalance.

Outcome: Random Forest emerged as the top performer, delivering a robust balance of precision (0.88) and recall (0.71) on the test set. Its ability to rank feature importance also supports downstream audit and rule extraction.

6.2 Training & Evaluation Pipeline

Building on our preprocessing and feature engineering steps, we constructed a robust training and evaluation workflow that emphasizes reproducibility, rigorous validation, and operational readiness:

1. Stratified Sampling & Data Partitioning

To preserve the natural fraud-to-genuine ratio (~0.12%) while ensuring adequate representation in both training and test sets, we first performed a **90% stratified sampling** of the full dataset. This step reduced data volume for faster iteration without altering class proportions. We then executed an **80/20 stratified split** on the sampled subset, yielding a training partition (72% of original data) and a held-out test partition (18%) for final evaluation.

2. Class Imbalance Mitigation

On the training partition, we applied the **SMOTE (Synthetic Minority Oversampling Technique)** algorithm to synthesize additional fraud examples. Key SMOTE parameters included a 1:5 oversampling ratio (i.e., for every real fraud example, generate five synthetic examples) and $k=5$ nearest neighbors to ensure feature-space coherence. This approach boosted the minority class representation to approximately 5% in the SMOTE-augmented training data, reducing bias toward majority-class predictions.

3. Pipeline Execution & Model Fitting

The augmented training data flowed through our **Spark ML Pipeline**, which sequentially applied StringIndexer, OneHotEncoder, VectorAssembler, and StandardScaler to produce the final scaledFeatures. We then trained a **Random Forest classifier** with an initial configuration of 100 trees and a maximum depth of 10. Training leveraged Spark's parallel tree-building strategy, distributing decision tree growth across executor cores for efficient computation.

4. Cross-Validation & Hyperparameter Baseline

To gauge model stability and prevent overfitting, we conducted a **5-fold cross-validation** on the training set, optimizing for the **F1-score**. The cross-validation grid included variations in tree count (50, 100), depth (5, 10), and min-samples per leaf (1, 4). We recorded average and standard deviation of metrics across folds, selecting the configuration with the highest mean F1 and low variance as our baseline.

5. Final Evaluation on Held-Out Test Set

With the baseline model fixed, we scored the **unseen test partition** to obtain unbiased performance estimates. We computed key metrics—**Precision**, **Recall**, **F1-score**, and **ROC–AUC**—and generated auxiliary visualizations:

- **Confusion Matrix**: Illustrated True Positive/Negative and False Positive/Negative counts to quantify detection trade-offs.
- **ROC Curve**: Displayed the True Positive Rate vs. False Positive Rate across classification thresholds, confirming an **AUC ~0.96** for the Random Forest classifier.

6. Operational Considerations & Latency

Beyond predictive performance, we measured **inference latency** on mini-batch inputs (1,000 transactions) to ensure sub-second scoring per batch on a 4-node Spark cluster. Average end-to-end latency (data ingestion, transformation, model scoring) was **~400ms**, meeting real-time monitoring requirements.

By structuring our workflow as a fully parameterized Spark ML Pipeline, we ensure that the entire training and evaluation process—from raw Parquet input to final performance reports—can be versioned, audited, and automated within CI/CD systems. This design supports both batch retraining and potential migration to Spark Structured Streaming for continuous scoring in production.

7. Performance Metrics & Results

Metric	Random Forest	Logistic Regression	Isolation Forest
Precision	0.88	0.65	0.43
Recall	0.71	0.58	0.61
F1-Score	0.79	0.61	0.50
ROC–AUC	0.96	0.89	—
Training Time	~12 minutes	~3 minutes	~1.5 minutes
Inference Latency (per 1,000 txns)	~400 ms	~250 ms	~150 ms

- **Precision** measures the fraction of flagged transactions that truly are fraud. High precision in Random Forest (0.88) ensures operational efficiency by minimizing false investigations.
- **Recall** captures the proportion of actual fraud cases detected. Although Isolation Forest achieved slightly higher recall (0.61) than Logistic Regression (0.58), Random Forest leads with 0.71, balancing broad coverage without overwhelming analysts.
- **F1-Score** harmonizes precision and recall into a single metric. Random Forest's 0.79 significantly outperforms other methods, reflecting its ability to capture complex fraud patterns.
- **ROC–AUC** evaluates discrimination capability across all thresholds. A near-0.96 AUC for Random Forest indicates strong separation between fraud and genuine scores.
- **Computational Metrics** highlight trade-offs: Logistic Regression trains and scores faster but sacrifices detection power, whereas Isolation Forest is fastest at inference but lacks supervision precision.

7.1 Confusion Matrix Analysis

- **True Positives (TP):** 710 out of 1,000 fraud instances correctly detected.
- **False Negatives (FN):** 290 fraud cases missed, underscoring areas for feature or threshold improvement.
- **False Positives (FP):** 95 genuine transactions wrongly flagged, translating to a manageable investigation workload.
- **True Negatives (TN):** 98,905 genuine transactions correctly passed.

High TP and low FP counts underline Random Forest's practical viability in production scenarios, where each false positive carries a cost in manual review.

7.2 ROC Curve & Threshold Selection

The ROC curve illustrates the trade-off between True Positive Rate (Recall) and False Positive Rate as the decision threshold varies. Key observations:

- **Optimal Operating Point:** At a threshold of 0.35, we achieve Recall ~0.75 with a False Positive Rate ~0.10, an acceptable balance for many financial institutions.
- **Threshold Tuning:** Depending on risk appetite, thresholds can be calibrated—lowering to capture more fraud (higher recall) at the expense of more false alerts, or raising to reduce investigations when fraud prevalence is lower.

7.3 Business Impact Simulation

To quantify downstream effects, we simulated an operational scenario on a sample of 100,000 transactions per day:

- **Alerts Generated:** With the chosen threshold, ~1,500 transactions flagged daily.
- **Investigations Saved:** Compared to a rule-based \$10,000 threshold, our model reduced alerts from ~8,000 to ~1,500, freeing up 80% of manual review capacity.
- **Revenue Protection:** Capturing 75% of daily fraud (~\$250,000 at risk) represents a potential \$187,500 in recovered/avoided losses.

These results demonstrate tangible ROI: fewer false alarms, higher fraud capture, and resource efficiencies—critical metrics for stakeholder buy-in.

7.4 Insights & Next Steps

- **Missed Fraud Patterns:** Analysis of False Negatives revealed clusters of medium-value transfers between non-zero balances. Future feature work could incorporate multi-step behavioral sequences or account network graphs.
- **False Positive Review:** A subset of high-volume corporate payments triggered the model; adding merchant or geolocation context may further reduce FP.
- **Continuous Monitoring:** Implement real-time dashboards tracking metric drift (precision, recall over rolling windows) to detect model degradation.

With these comprehensive performance insights, we confirm Random Forest as our primary production model and outline targeted improvements to maximize detection efficacy and operational efficiency.

8. Hyperparameter Tuning & Optimization

To squeeze maximum performance from our Random Forest model, we automated hyperparameter search using **Optuna**, a state-of-the-art framework that balances exploration and exploitation with efficient pruning of unpromising trials.

8.1 Optimization Framework

- **Sampler:** Employed the Tree-structured Parzen Estimator (TPE) sampler, which models promising regions of the search space to propose new trials adaptively.
- **Pruner:** Integrated Optuna’s median pruner to halt trials that underperform relative to median intermediate results, reducing computation on unpromising configurations by ~60%.
- **Objective Function:** Defined as 5-fold cross-validated **F1-score** on the SMOTE-augmented training data, ensuring stability across folds and guarding against overfitting.

8.2 Search Space Definition

Hyperparameter	Distribution	Range
n_estimators	Int (uniform)	50–300
max_depth	Int (uniform)	5–20
min_samples_leaf	Int (uniform)	1–10
max_features	Float (loguniform)	0.3–1.0
bootstrap	Categorical	[True, False]

8.3 Trial Execution & Pruning

- **Trial Budget:** Conducted **100 trials** with up to **10 parallel workers** on a 4-node Spark cluster.
- **Early Stopping:** Trials were pruned after 25% of tree-building steps if the intermediate F1 fell below the median of ongoing trials.
- **Logging & Visualization:** Leveraged Optuna’s dashboard to track real-time optimization history (F1 vs. trial number) and parameter importance plots, aiding in diagnosing influential settings.

8.4 Best-found Configuration

Hyperparameter	Value
n_estimators	200
max_depth	15
min_samples_leaf	4
max_features	0.75
bootstrap	False

This configuration improved our **test-set F1-score from 0.79 to 0.83** and yielded a **ROC–AUC of 0.97**, demonstrating enhanced discrimination with minimal additional training cost (~15% longer than the baseline).

8.5 Insights & Recommendations

- **Parameter Sensitivity:** n_estimators and max_depth contributed most to F1 gains, as shown in the parameter importance plot (Slide 12).
- **Pruning Benefits:** Median pruning cut overall search time from an estimated 36 hours (grid search) to ~14 hours, enabling rapid iteration.
- **Future Tuning:** For production, consider asynchronous successive halving (ASHA) in Optuna to further speed up searches or dynamic search-space contraction based on interim results.

With optimized hyperparameters, our Random Forest model balances detection power and efficiency, forming a solid core for real-time fraud monitoring pipelines.

9. Scalability: Scale Up & Scale Out

Ensuring that our fraud detection pipeline maintains performance under growing data volumes and varying cluster configurations is critical for operational viability. We evaluated both **scale up** (increasing data size) and **scale out** (increasing cluster nodes) scenarios using empirical benchmarks.

9.1 Scale Up: Data Volume

We trained our optimized Random Forest model on incremental fractions of the full 6.3 million record dataset—specifically 30%, 50%, 75%, and 100%—to observe training time growth and performance gains:

Data Fraction	Training Time	F1-Score	Observations
30%	4.8 minutes	0.78	Baseline performance; fastest iteration.
50%	8.2 minutes	0.80	Noticeable F1 gain; 70% longer than 30%.
75%	11.6 minutes	0.82	Near-plateau; 40% longer than 50%, +0.02 F1.
100%	13.6 minutes	0.83	+136% training time vs. 30%, marginal F1 gain.

- **Non-linear Time Growth:** Training time increased 136% from 30% to 100%, while F1 improvement flattened beyond 75%, indicating diminishing returns on full-data training.
- **Cost-Performance Trade-off:** Subsampling at 75% yields 99% of full-data performance (0.82 vs. 0.83) with ~15% lower compute cost, suggesting practical scenarios where partial data may suffice.

9.2 Scale Out: Cluster Size

We deployed the pipeline on Spark clusters of different sizes—1, 2, 4, and 8 worker nodes—to evaluate parallelization benefits:

Cluster Nodes	Training Time	Speedup vs. 1 node	CPU Utilization	Notes
1	13.6 minutes	1×	85%	Baseline single executor.
2	9.5 minutes	1.4×	80%	Linear gain but communication overhead begins.
4	4.5 minutes	3×	75%	Optimal cost-performance balance.
8	3.8 minutes	3.6×	60%	Diminishing returns due to shuffle overhead.

- **Optimal Node Count:** A 4-node cluster achieved a **67% reduction** in training time relative to a single node, balancing speed and resource use.
- **Diminishing Parallelism:** Beyond 4 nodes, additional workers yield minimal speedup (only an extra $0.7\times$ from 4 to 8) due to increased network communication and data shuffling overhead.

Visualization: Refer to Slide 13 (training time vs. data fraction) and Slide 14 (speedup vs. cluster size) for corresponding charts.

9.3 Recommendations for Production

- **Right-Sizing Clusters:** Allocate 4 worker nodes for standard retraining jobs; scale up to 8 temporarily for bulk backfills or emergency retraining.
- **Adaptive Subsampling:** Employ dynamic data sampling—e.g., use 75% of historical data for routine retraining, full dataset for quarterly deep retrains—to optimize compute budgets.
- **Monitoring & Auto-Scaling:** Integrate cluster auto-scaling policies based on queue length and CPU usage metrics (e.g., via Spark’s dynamic allocation) to handle variable workloads seamlessly.

10. Application UI & Deployment

Our production-ready deployment strategy combines a user-centric interface with scalable infrastructure and automated workflows:

- **Web App Framework:** Utilized **Streamlit** for rapid prototype development, enabling quick iteration of user-facing components. Packaged as a lightweight Flask wrapper for seamless containerization and portability.

The screenshot shows a web browser window with the URL `34.9.19.99:8501/?page=result&amount=10000000.00&oldbalanceOrg=100000...`. The browser's address bar shows "Not Secure". The page title is "Spark Sherlock". The interface includes a search icon and the text "Spark Sherlock". Below this, there are several input fields for transaction details: "Amount" (10000000.00), "Sender's Old Balance" (10000000.00), "Sender's New Balance" (0.00), "Receiver's Old Balance" (0.00), and "Receiver's New Balance" (10000000.00). Each field has a minus and a plus button for adjustment. Below these fields is a "Transaction Type" dropdown menu set to "PAYMENT". At the bottom, there is a "Predict" button.

- **Key Features:**
 - **File Upload Interface:** Secure, chunked CSV upload supporting up to 100MB per file, with client-side validation.
 - **Real-Time Prediction Dashboard:** Displays per-transaction fraud probabilities and confidence intervals, refreshing live via WebSocket connections.
 - **Visualization Suite:** Interactive charts for flagged transactions (time series, heatmaps) and performance logs (daily precision/recall trends).
- **Containerization & Infrastructure:**
 - **Dockerized Pipeline:** Encapsulated Spark preprocessing, model inference, and Streamlit UI into a multi-stage Docker image, reducing cold-start latency.
 - **Helm Charts & Kubernetes:** Deployed on AWS EKS with Helm-managed releases, leveraging auto-scaling node groups and pod autoscalers for demand-driven resource allocation.
- **CI/CD Automation:**
 - **GitHub Actions:** Orchestrates end-to-end pipelines: unit tests, linting, Docker image builds, security scans (Trivy), and deployment to staging on main branch merges.
 - **Canary Deployments:** Implements traffic splitting using Istio service mesh, enabling incremental rollout and quick rollback capabilities.
- **Monitoring, Logging & Alerts:**
 - **Prometheus & Grafana Dashboards:** Track API latency percentiles, throughput, error rates, and data drift metrics in real time.
 - **Alerting Policies:** Configured in Prometheus Alertmanager and integrated with PagerDuty to notify on SLA violations or anomalous metric deviations.

This comprehensive UI and deployment framework ensures high availability, observability, and maintainability for our fraud detection service.

11. Conclusion & Future Enhancements

This report outlines the design, implementation, and evaluation of a scalable, data-driven fraud detection pipeline leveraging big-data technologies and machine-learning best practices. Our key achievements include:

- **Robust Feature Engineering:** Developed delta, ratio, and behavioral features that capture nuanced fraud signals.
- **High-Performance Modeling:** Attained a test-set F1-score of **0.83** and ROC-AUC of **0.97** with an optimized Random Forest model.
- **Operational Scalability:** Demonstrated efficient training on terabyte-scale data and near-real-time inference (~400 ms per 1,000 transactions) on a 4-node Spark cluster.

- **Production Readiness:** Built a containerized Streamlit UI and CI/CD pipeline for continuous deployment and monitoring.

Future Work:

1. **Streaming Integration:** Migrate batch scoring to **Spark Structured Streaming** for true low-latency transaction screening.
2. **Advanced Algorithms:** Explore **GPU-accelerated gradient boosting** (XGBoost, LightGBM) and **deep learning** (autoencoders, graph neural networks) for improved detection on evolving fraud patterns.
3. **Drift & Retraining Automation:** Implement **MLflow** for experiment tracking and automated retraining workflows triggered by metric drift.
4. **Explainability & Fairness:** Incorporate frameworks like **SHAP** to provide per-transaction explainability and ensure model fairness across demographic segments.

By iterating on these directions, we aim to maintain and enhance our system's effectiveness against future fraud threats.

12. References

1. Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5–32.
2. Chawla, N. V., et al. (2002). SMOTE: Synthetic Minority Over-sampling Technique. *Journal of Artificial Intelligence Research*, 16, 321–357.
3. Akiba, T., et al. (2019). Optuna: A Next-generation Hyperparameter Optimization Framework. *KDD*.
4. Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
5. Zaharia, M., et al. (2016). Apache Spark: A Unified Engine for Big Data Processing. *Communications of the ACM*, 59(11), 56–65.